

Efficient renaming in Conflict-free Replicated Data Types (CRDTs)

Matthieu Nicolas

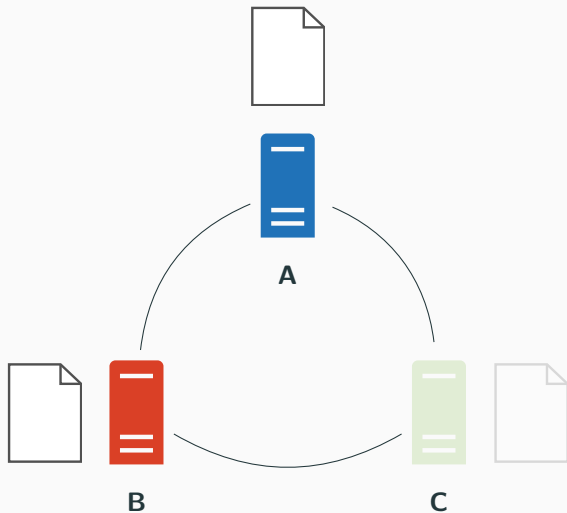
COAST team

Supervised by Gérald Oster and Olivier Perrin

December 4, 2018

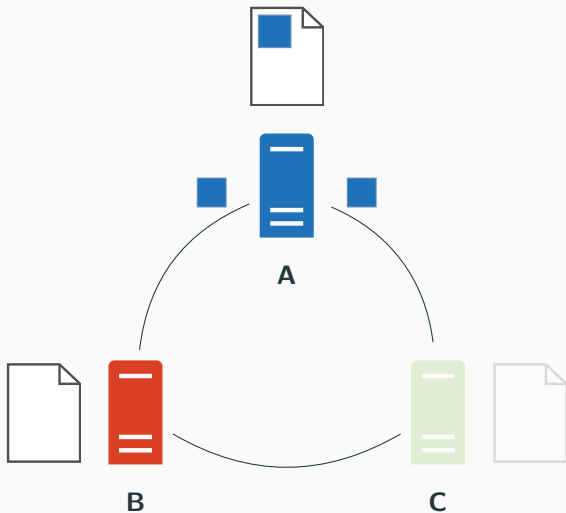


Conflict-free Replicated Data Types (CRDTs) [3]



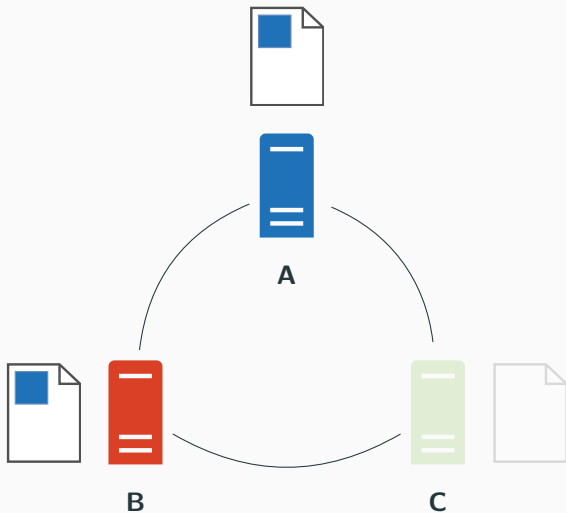
- Replicated data structure

Conflict-free Replicated Data Types (CRDTs) [3]



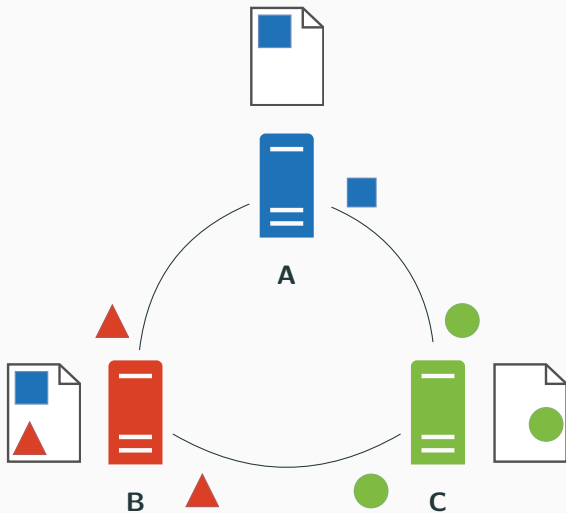
- Replicated data structure
- Updates performed without coordination

Conflict-free Replicated Data Types (CRDTs) [3]



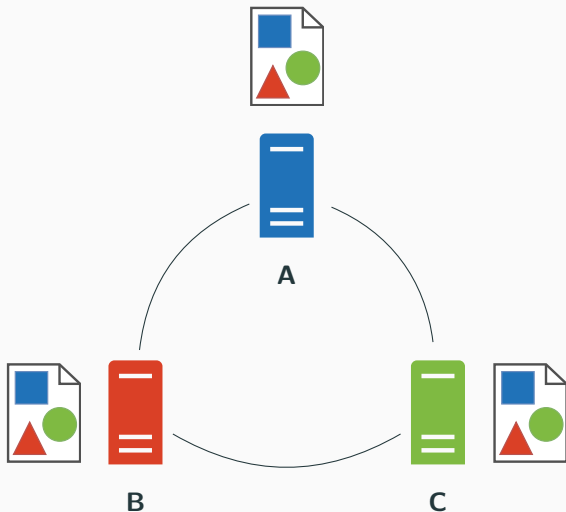
- Replicated data structure
- Updates performed without coordination

Conflict-free Replicated Data Types (CRDTs) [3]



- Replicated data structure
- Updates performed without coordination

Conflict-free Replicated Data Types (CRDTs) [3]



- Replicated data structure
- Updates performed without coordination
- Strong Eventual Consistency [3]

Identifier-based CRDTs

Main idea

- Attach an identifier to each element

Allow to design commutative updates

- Identifying uniquely elements
- Ordering updates causally
- ...

Limits

- Unbounded size of identifiers
- Overhead of the data structure increasing over time

How to reduce the overhead introduced by the data structure ?

Proposition

- Reassign shorter identifiers in a fully distributed manner
- Focus on one CRDT

- State of the art of *Sequence CRDTs*
- Elements are ordered by their identifier, noted here as lowercase letters

- State of the art of *Sequence CRDTs*
- Elements are ordered by their identifier, noted here as lowercase letters



Figure 1: State of a sequence which contains the elements "helo" and their corresponding identifiers

- State of the art of *Sequence CRDTs*
- Elements are ordered by their identifier, noted here as lowercase letters



Figure 1: State of a sequence which contains the elements "helo" and their corresponding identifiers



Figure 2: State of a sequence which contains the block "helo"

Example




Figure 3: Example of concurrent *insert* operations

Example




Figure 3: Example of concurrent *insert* operations

Example




Figure 3: Example of concurrent *insert* operations

Example




Figure 3: Example of concurrent *insert* operations

Example




Figure 3: Example of concurrent *insert* operations

Declining performances

- Updates performed may lead to an inefficient internal representation




Figure 4: Example of inefficient internal representation

- The more blocks we have:
 - The more metadata we store
 - The longer it takes to browse the sequence to *insert* or *delete* an element

Renaming mechanism

- Introduce a *rename* operation




Figure 5: Example of renaming

Renaming mechanism

- Introduce a *rename* operation




Figure 5: Example of renaming

- Generates a new identifier to the first element, based on its previous identifier

Renaming mechanism

- Introduce a *rename* operation



Figure 5: Example of renaming

- Generates a new identifier to the first element, based on its previous identifier
- Then generates contiguous identifiers for all following elements

Renaming mechanism

- Introduce a *rename* operation



Figure 5: Example of renaming

- Generates a new identifier to the first element, based on its previous identifier
- Then generates contiguous identifiers for all following elements

Renaming mechanism

- Introduce a *rename* operation




Figure 5: Example of renaming

- Generates a new identifier to the first element, based on its previous identifier
- Then generates contiguous identifiers for all following elements

Handling concurrent operations

- Others may perform updates concurrently to a *rename* operation




Figure 6: Example of concurrent insert

Handling concurrent operations

- Others may perform updates concurrently to a *rename* operation




Figure 6: Example of concurrent insert

Handling concurrent operations

- Others may perform updates concurrently to a *rename* operation




Figure 6: Example of concurrent insert

Handling concurrent operations

- Others may perform updates concurrently to a *rename* operation
- May lead to inconsistencies




Figure 6: Example of concurrent insert

Handling concurrent operations

- Others may perform updates concurrently to a *rename* operation
- May lead to inconsistencies




Figure 6: Example of concurrent insert

- Define rewriting rules to transform identifiers from one *epoch* to another

Handling concurrent operations

- Others may perform updates concurrently to a *rename* operation
- May lead to inconsistencies




Figure 6: Example of concurrent insert

- Define rewriting rules to transform identifiers from one *epoch* to another

Handling concurrent rename

- Define a total order between *rename* operations
- Pick a "winner" operation between concurrent *renames*
- Define additional rewriting rules to *undo* the effect of "losing" ones

To wrap up

Done

- Designed a *rename* operation for LogootSplit
- Defined rewriting rules to deal with concurrent updates

Work in progress

- Implementation in MUTE [2], our P2P collaborative text editor
- Design the strategy to trigger the renaming

To do

- Prove formally the correctness of the mechanism
- Benchmark its performances

Generalize the approach

- To other Sequence CRDTs
- To other types
 - Counter
 - Set
 - ...

Thanks for your attention, any questions?



- [1] L. André, S. Martin, G. Oster, and C.-L. Ignat.

Supporting adaptable granularity of changes for massive-scale collaborative editing.

In International Conference on Collaborative Computing: Networking, Applications and Worksharing - CollaborateCom 2013, pages 50–59, Austin, TX, USA, Oct. 2013. IEEE Computer Society.

- [2] M. Nicolas, V. Elvinger, G. Oster, C.-L. Ignat, and F. Charoy.

MUTE: A Peer-to-Peer Web-based Real-time Collaborative Editor.

In ECSCW 2017 - 15th European Conference on Computer-Supported Cooperative Work, volume 1 of *Proceedings of 15th European Conference on Computer-Supported Cooperative Work - Panels, Posters and Demos*, pages 1–4, Sheffield, United Kingdom, Aug. 2017. EUSSET.

- [3] M. Shapiro, N. M. Preguiça, C. Baquero, and M. Zawirski.
Conflict-free replicated data types.
In *Proceedings of the 13th International Symposium on Stabilization, Safety, and Security of Distributed Systems*, SSS 2011, pages 386–400, 2011.