

# Efficient renaming in Conflict-free Replicated Data Types (CRDTs)

---

Matthieu Nicolas

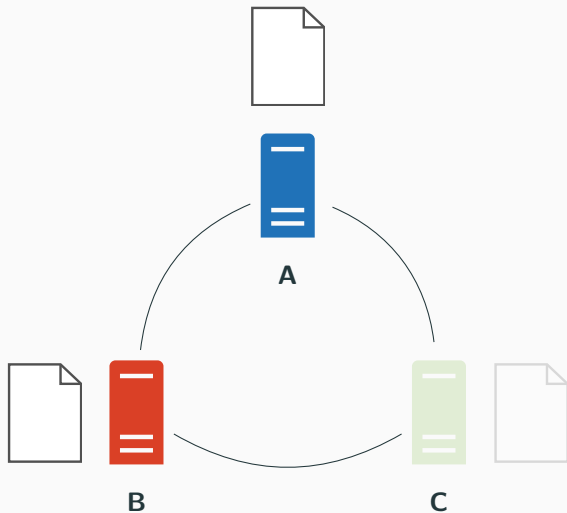
COAST team

**Supervised by** Gérald Oster and Olivier Perrin

November 29, 2018

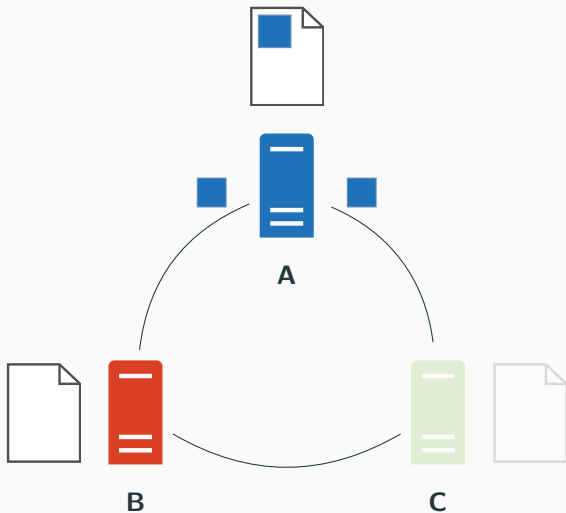


# Conflict-free Replicated Data Types (CRDTs) [3]



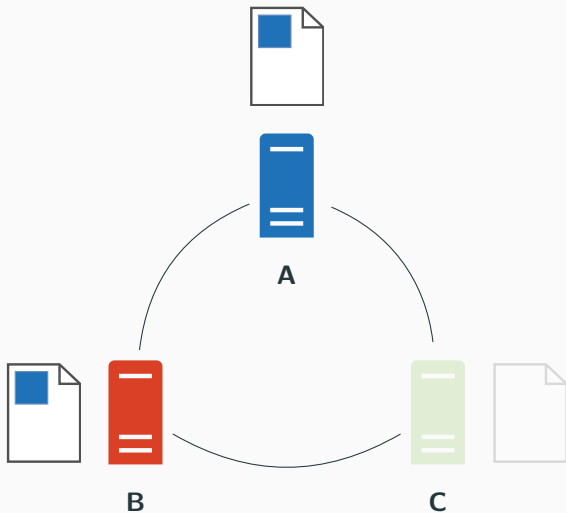
- Replicated data structure

# Conflict-free Replicated Data Types (CRDTs) [3]



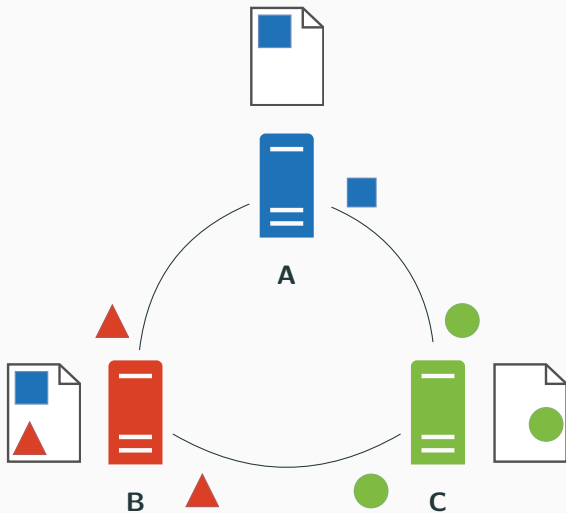
- Replicated data structure
- Updates performed without coordination

# Conflict-free Replicated Data Types (CRDTs) [3]



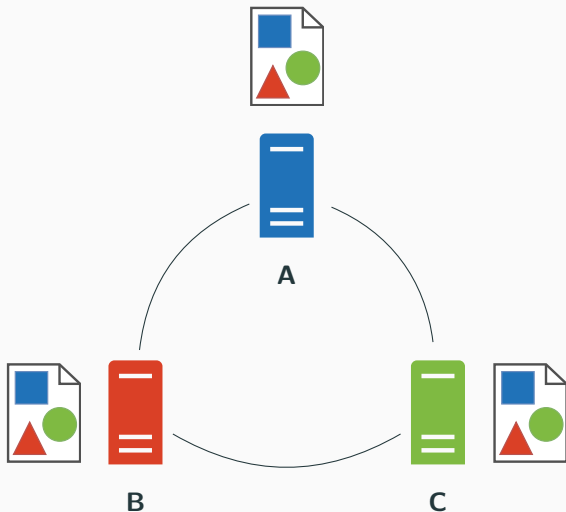
- Replicated data structure
- Updates performed without coordination

# Conflict-free Replicated Data Types (CRDTs) [3]



- Replicated data structure
- Updates performed without coordination

# Conflict-free Replicated Data Types (CRDTs) [3]



- Replicated data structure
- Updates performed without coordination
- Strong Eventual Consistency [3]

# Identifier-based CRDTs

## Main idea

- Attach an identifier to each inserted element

## Allow to achieve commutative updates

- By identifying uniquely elements
- By ordering them relatively to each other

- State of the art of *Sequence CRDTs*
- Relies on *identifiers* to ensure convergence, noted here as letters and words
- Elements are ordered by their identifier



- State of the art of *Sequence CRDTs*
- Relies on *identifiers* to ensure convergence, noted here as letters and words
- Elements are ordered by their identifier



**Figure 1:** The state of a sequence which contains the elements "helo" and their corresponding identifiers

# LogootSplit [1]

- State of the art of *Sequence CRDTs*
- Relies on *identifiers* to ensure convergence, noted here as letters and words
- Elements are ordered by their identifier

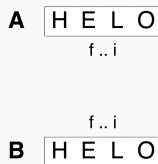


**Figure 1:** The state of a sequence which contains the elements "helo" and their corresponding identifiers




**Figure 2:** The state of a sequence which contains the block "helo"

# Example



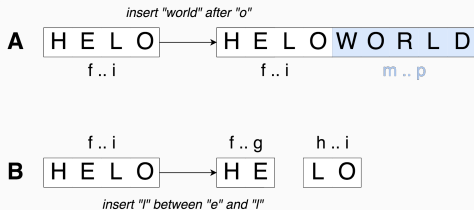
**Figure 3:** Example of concurrent *insert* operations

# Example




**Figure 3:** Example of concurrent *insert* operations

# Example



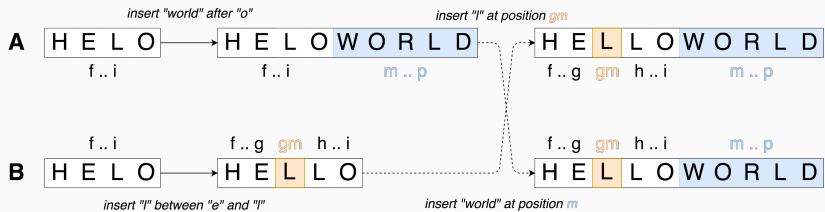
**Figure 3:** Example of concurrent *insert* operations

# Example



**Figure 3:** Example of concurrent *insert* operations

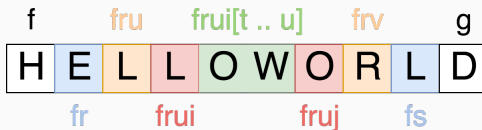
# Example



**Figure 3:** Example of concurrent *insert* operations

# Declining performances

- Operations performed may lead to an inefficient internal representation



**Figure 4:** Example of inefficient internal representation


- The more blocks we have:
  - The more metadata we store
  - The longer it takes to browse the sequence to *insert* or *delete* an element



**How to reduce the footprint of the metadata ?**

# Renaming mechanism

- Introduce a *rename* operation



**Figure 5:** Example of renaming

# Renaming mechanism

- Introduce a *rename* operation




**Figure 5:** Example of renaming

- Generates a new identifier to the first element, based on its previous identifier

# Renaming mechanism

- Introduce a *rename* operation



**Figure 5:** Example of renaming

- Generates a new identifier to the first element, based on its previous identifier
- Then generates contiguous identifiers for all following elements

# Renaming mechanism

- Introduce a *rename* operation




**Figure 5:** Example of renaming

- Generates a new identifier to the first element, based on its previous identifier
- Then generates contiguous identifiers for all following elements

# Renaming mechanism

- Introduce a *rename* operation




**Figure 5:** Example of renaming

- Generates a new identifier to the first element, based on its previous identifier
- Then generates contiguous identifiers for all following elements

# Renaming mechanism

- Introduce a *rename* operation




**Figure 5:** Example of renaming

- Generates a new identifier to the first element, based on its previous identifier
- Then generates contiguous identifiers for all following elements
- Each *rename* marks the beginning of a new *epoch*

# Handling concurrent operations

- Others can perform updates concurrently to a *rename* operation
- Apply them as such could lead to inconsistencies




**Figure 6:** Example of inconsistency



# Rewriting concurrent operations

- Define rewriting rules to transform identifiers from one *epoch* to another
- Upon reception of concurrent operations, rewrite identifiers before applying them



**Figure 7:** Example of rewriting

# Handling concurrent rename

- Define a total order between *rename* operations
- Pick a "winner" operation between concurrent *renames*
- Define additional rewriting rules to *undo* the effect of "losing" ones

- Propose a fully distributed renaming mechanism for *LogootSplit*
- Allows to reinitialize the footprint of the CRDT without coordination
- Implemented in MUTE [2], our P2P collaborative text editor

# Next steps

## **Provide a formal proof**

- Need to ensure the correctness of our algorithm

## **Benchmark the mechanism**

- Measure its impact on the performances
- Compare different strategies

## **Generalize the approach to other CRDTs**

Thanks for your attention, any questions?



- [1] L. André, S. Martin, G. Oster, and C.-L. Ignat.

**Supporting adaptable granularity of changes for massive-scale collaborative editing.**

*In International Conference on Collaborative Computing: Networking, Applications and Worksharing - CollaborateCom 2013*, pages 50–59, Austin, TX, USA, Oct. 2013. IEEE Computer Society.

- [2] M. Nicolas, V. Elvinger, G. Oster, C.-L. Ignat, and F. Charoy.

**MUTE: A Peer-to-Peer Web-based Real-time Collaborative Editor.**

*In ECSCW 2017 - 15th European Conference on Computer-Supported Cooperative Work*, volume 1 of *Proceedings of 15th European Conference on Computer-Supported Cooperative Work - Panels, Posters and Demos*, pages 1–4, Sheffield, United Kingdom, Aug. 2017. EUSSET.

- [3] M. Shapiro, N. M. Preguiça, C. Baquero, and M. Zawirski.  
**Conflict-free replicated data types.**  
In *Proceedings of the 13th International Symposium on Stabilization, Safety, and Security of Distributed Systems, SSS 2011*, pages 386–400, 2011.